



Schneider Electric Machine Learning Project

PRASAD JADHAV

```
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv('schneider_electric_dataset.csv')
pd.set_option('display.max_columns', 20)
print(df.shape)

(100000, 20)

df.head()
```

| | Product_ID | Category | Technical_Specifications | Sales |
|----------|-----------------|-------------|-------------------------------------|---------|
| Region \ | | | | |
| 0 North | 1 | Power Meter | {"voltage": "220V", "power": "50W"} | 4851 |
| 1 West | 2 | Switch | {"voltage": "220V", "power": "50W"} | 556 |
| 2 East | 3 | Switch | {"voltage": "110V", "power": "30W"} | 4713 |
| 3 East | 4 | Relay | {"voltage": "220V", "power": "50W"} | 4499 |
| 4 West | 5 | Switch | {"voltage": "110V", "power": "30W"} | 3704 |
| | | | | |
| | Customer Rating | Fault Flag | Maintenance Records | Price \ |

| | | | | |
|---|------|-----|---|---------|
| 0 | 3.62 | No | 5 | 2009.28 |
| 1 | 1.70 | No | 2 | 1899.78 |
| 2 | 1.80 | Yes | 0 | 2271.00 |
| 3 | 3.46 | No | 2 | 1174.32 |
| 4 | 4.59 | No | 7 | 3081.87 |

| | Date | Energy_Consumption \ |
|---|---------------------|----------------------|
| 0 | 2020-01-01 00:00:00 | 468.20 |
| 1 | 2020-01-01 01:00:00 | 133.46 |
| 2 | 2020-01-01 02:00:00 | 414.33 |
| 3 | 2020-01-01 03:00:00 | 170.97 |
| 4 | 2020-01-01 04:00:00 | 115.83 |

| | Sensor_Readings | Warranty_Expiry | Lead_Time \ |
|---|--------------------------------------|-----------------|-------------|
| 0 | {"temperature": 30, "pressure": 1.2} | 24 | 11 |
| 1 | {"temperature": 40, "pressure": 1.5} | 9 | 15 |
| 2 | {"temperature": 40, "pressure": 1.5} | 7 | 10 |
| 3 | {"temperature": 40, "pressure": 1.5} | 26 | 18 |
| 4 | {"temperature": 40, "pressure": 1.5} | 30 | 19 |

| | Supply_Chain_Risk | Demand_Variation | Promotions_Applied |
|---------|--------------------|------------------|--------------------|
| 0 | Competitor_Price \ | | |
| | 0.13 | 45.39 | Yes |
| 2990.56 | | | |
| 1 | 0.40 | 47.75 | Yes |
| 1358.28 | | | |
| 2 | 0.94 | 46.11 | No |
| 2160.00 | | | |
| 3 | 0.08 | 11.98 | No |
| 2948.44 | | | |
| 4 | 0.06 | 35.71 | No |
| 4197.57 | | | |

| | Return_Flag | User_Demographics |
|---|-------------|-------------------|
| 0 | Yes | Enterprise |
| 1 | Yes | Individual |
| 2 | No | Small Business |
| 3 | Yes | Government |
| 4 | Yes | Individual |

df.tail()

| | Product_ID | Category | Technical_Specifications | Sales |
|----------|------------|----------|-------------------------------------|-------|
| Region \ | | | | |
| 99995 | 99996 | Switch | {"voltage": "220V", "power": "50W"} | 3311 |
| West | | | | |
| 99996 | 99997 | Switch | {"voltage": "220V", "power": "50W"} | 4178 |
| East | | | | |
| 99997 | 99998 | Switch | {"voltage": "220V", "power": "50W"} | 4403 |
| West | | | | |

| | | | | |
|-------------|--------------------------------------|--------------------|-------------------------------------|---------|
| 99998 | 99999 | Switch | {"voltage": "110V", "power": "30W"} | 2154 |
| West | | | | |
| 99999 | 100000 | Switch | {"voltage": "220V", "power": "50W"} | 4157 |
| North | | | | |
| | Customer_Rating | Fault_Flag | Maintenance_Records | Price \ |
| 99995 | 2.41 | Yes | 6 | 4651.82 |
| 99996 | 4.82 | No | 2 | 3206.15 |
| 99997 | 1.25 | Yes | 9 | 280.02 |
| 99998 | 4.65 | No | 7 | 4570.44 |
| 99999 | 1.58 | Yes | 8 | 2879.91 |
| | Date | Energy_Consumption | \ | |
| 99995 | 2031-05-29 11:00:00 | 445.16 | | |
| 99996 | 2031-05-29 12:00:00 | 292.51 | | |
| 99997 | 2031-05-29 13:00:00 | 300.55 | | |
| 99998 | 2031-05-29 14:00:00 | 279.96 | | |
| 99999 | 2031-05-29 15:00:00 | 382.38 | | |
| | | Sensor_Readings | Warranty_Expiry | |
| Lead_Time \ | | | | |
| 99995 | {"temperature": 30, "pressure": 1.2} | | | 6 |
| 19 | | | | |
| 99996 | {"temperature": 30, "pressure": 1.2} | | | 3 |
| 11 | | | | |
| 99997 | {"temperature": 30, "pressure": 1.2} | | | 9 |
| 5 | | | | |
| 99998 | {"temperature": 30, "pressure": 1.2} | | | 16 |
| 2 | | | | |
| 99999 | {"temperature": 40, "pressure": 1.5} | | | 3 |
| 11 | | | | |
| | Supply_Chain_Risk | Demand_Variation | Promotions_Applied | \ |
| 99995 | 0.16 | 47.12 | Yes | |
| 99996 | 0.70 | 32.99 | Yes | |
| 99997 | 0.79 | 4.98 | Yes | |
| 99998 | 0.23 | 23.90 | Yes | |
| 99999 | 0.21 | 18.29 | No | |
| | Competitor_Price | Return_Flag | User_Demographics | |
| 99995 | 543.99 | Yes | Small Business | |
| 99996 | 4132.05 | Yes | Enterprise | |
| 99997 | 2179.27 | No | Government | |
| 99998 | 2823.15 | No | Government | |
| 99999 | 1819.78 | Yes | Enterprise | |

Dataset Columns Overview

- Product_ID: Unique identifier for each product.
- Category: Product category (e.g., circuit breakers, switches).

- Technical_Specifications: JSON field containing specs (voltage, power rating, etc.).
- Sales: Historical sales data (number of units sold).
- Region: Sales region.
- Customer_Rating: Average customer ratings.
- Fault_Flag: Indicates if a product had a fault (Yes/No).
- Maintenance_Records: Number of maintenance events.
- Price: Product price.
- Date: Timestamp of sales/maintenance data.
- Energy_Consumption: Energy consumed by the product (kWh).
- Sensor_Readings: Data from product sensors (temperature, pressure, etc.).
- Warranty_Expiry: Time remaining on the product warranty (in months).
- Lead_Time: Delivery lead time (in days).
- Supply_Chain_Risk: Risk factor based on supply chain conditions.
- Demand_Variation: Month-to-month variation in demand.
- Promotions_Applied: Whether a promotion was applied (Yes/No).
- Competitor_Price: Price of a similar product by competitors.
- Return_Flag: Indicates if the product was returned (Yes/No).
- User_Demographics: Information about the purchasing company or customer (size, industry).

Data Exploration

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 100000 entries, 0 to 99999
```

```
Data columns (total 20 columns):
```

| # | Column | Non-Null Count | Dtype |
|----|--------------------------|-----------------|---------|
| 0 | Product_ID | 100000 non-null | int64 |
| 1 | Category | 100000 non-null | object |
| 2 | Technical_Specifications | 100000 non-null | object |
| 3 | Sales | 100000 non-null | int64 |
| 4 | Region | 100000 non-null | object |
| 5 | Customer_Rating | 100000 non-null | float64 |
| 6 | Fault_Flag | 100000 non-null | object |
| 7 | Maintenance_Records | 100000 non-null | int64 |
| 8 | Price | 100000 non-null | float64 |
| 9 | Date | 100000 non-null | object |
| 10 | Energy_Consumption | 100000 non-null | float64 |
| 11 | Sensor_Readings | 100000 non-null | object |
| 12 | Warranty_Expiry | 100000 non-null | int64 |
| 13 | Lead_Time | 100000 non-null | int64 |
| 14 | Supply_Chain_Risk | 100000 non-null | float64 |
| 15 | Demand_Variation | 100000 non-null | float64 |
| 16 | Promotions_Applied | 100000 non-null | object |
| 17 | Competitor_Price | 100000 non-null | float64 |

```
18 Return_Flag          100000 non-null object
19 User_Demographics    100000 non-null object
dtypes: float64(6), int64(5), object(9)
memory usage: 15.3+ MB
```

```
df.isnull().sum()
```

```
Product_ID          0
Category            0
Technical_Specifications  0
Sales              0
Region             0
Customer_Rating     0
Fault_Flag         0
Maintenance_Records 0
Price              0
Date              0
Energy_Consumption 0
Sensor_Readings     0
Warranty_Expiry     0
Lead_Time          0
Supply_Chain_Risk   0
Demand_Variation    0
Promotions_Applied  0
Competitor_Price    0
Return_Flag         0
User_Demographics   0
dtype: int64
```

```
df.duplicated().sum()
```

```
0
```

```
df.describe()
```

| | Product_ID | Sales | Customer_Rating |
|-----------------------|---------------|---------------|-----------------|
| Maintenance_Records \ | | | |
| count | 100000.000000 | 100000.000000 | 100000.000000 |
| mean | 50000.500000 | 2519.835300 | 3.008512 |
| std | 28867.657797 | 1430.183332 | 1.153912 |
| min | 1.000000 | 50.000000 | 1.000000 |
| 25% | 25000.750000 | 1280.000000 | 2.010000 |
| 50% | 50000.500000 | 2514.000000 | 3.010000 |
| 75% | 75000.250000 | 3766.000000 | 4.010000 |

| | | | |
|---------------|---------------|--------------------|-----------------|
| max | 100000.000000 | 4999.000000 | 5.000000 |
| 9.000000 | | | |
| | Price | Energy_Consumption | Warranty_Expiry |
| Lead_Time \ | | | |
| count | 100000.000000 | 100000.000000 | 100000.000000 |
| 100000.000000 | | | |
| mean | 2525.309220 | 254.743850 | 17.969350 |
| 9.991390 | | | |
| std | 1427.492425 | 141.423348 | 10.084338 |
| 5.478388 | | | |
| min | 50.010000 | 10.000000 | 1.000000 |
| 1.000000 | | | |
| 25% | 1289.540000 | 132.150000 | 9.000000 |
| 5.000000 | | | |
| 50% | 2527.080000 | 255.325000 | 18.000000 |
| 10.000000 | | | |
| 75% | 3760.700000 | 376.580000 | 27.000000 |
| 15.000000 | | | |
| max | 4999.870000 | 500.000000 | 35.000000 |
| 19.000000 | | | |

| | | | |
|-------|-------------------|------------------|------------------|
| | Supply_Chain_Risk | Demand_Variation | Competitor_Price |
| count | 100000.000000 | 100000.000000 | 100000.000000 |
| mean | 0.499296 | 25.093246 | 2558.876172 |
| std | 0.288830 | 14.402074 | 1416.896620 |
| min | 0.000000 | 0.000000 | 100.000000 |
| 25% | 0.250000 | 12.640000 | 1332.490000 |
| 50% | 0.500000 | 25.140000 | 2560.285000 |
| 75% | 0.750000 | 37.592500 | 3789.760000 |
| max | 1.000000 | 50.000000 | 4999.900000 |

```
num_cols = [x for x in df.columns if df[x].dtypes != 'float64']
```

```
for col in num_cols:
    print(f"Value counts for column '{col}':")
    print(df[col].value_counts())
    print("\n" + "_"*40 + "\n")
```

Value counts for column 'Product_ID':

| | |
|------------|---|
| Product_ID | |
| 1 | 1 |
| 66651 | 1 |
| 66673 | 1 |
| 66672 | 1 |
| 66671 | 1 |
| .. | |
| 33332 | 1 |
| 33331 | 1 |
| 33330 | 1 |

```
33329      1
100000     1
Name: count, Length: 100000, dtype: int64
```

Value counts for column 'Category':

```
Category
Circuit Breaker    20166
Switch             20024
Busway             20006
Power Meter        19945
Relay              19859
Name: count, dtype: int64
```

Value counts for column 'Technical_Specifications':

```
Technical_Specifications
{"voltage": "220V", "power": "50W"}    50074
{"voltage": "110V", "power": "30W"}    49926
Name: count, dtype: int64
```

Value counts for column 'Sales':

```
Sales
3333      39
1087      37
1437      37
4053      36
3482      35
..
3437       8
3549       8
3773       7
1925       7
1113       6
Name: count, Length: 4950, dtype: int64
```

Value counts for column 'Region':

```
Region
East      25177
South     25049
North     24887
West      24887
Name: count, dtype: int64
```

Value counts for column 'Fault_Flag':

Fault_Flag

Yes 50217

No 49783

Name: count, dtype: int64

Value counts for column 'Maintenance_Records':

Maintenance_Records

0 10146

8 10103

2 10071

9 10071

4 9988

1 9972

5 9968

6 9960

7 9939

3 9782

Name: count, dtype: int64

Value counts for column 'Date':

Date

2020-01-01 00:00:00 1

2027-08-09 02:00:00 1

2027-08-10 00:00:00 1

2027-08-09 23:00:00 1

2027-08-09 22:00:00 1

..

2023-10-20 19:00:00 1

2023-10-20 18:00:00 1

2023-10-20 17:00:00 1

2023-10-20 16:00:00 1

2031-05-29 15:00:00 1

Name: count, Length: 100000, dtype: int64

Value counts for column 'Sensor_Readings':

Sensor_Readings

{"temperature": 40, "pressure": 1.5} 50084

{"temperature": 30, "pressure": 1.2} 49916

Name: count, dtype: int64

Value counts for column 'Warranty_Expiry':

Warranty_Expiry

| | |
|----|------|
| 13 | 3026 |
| 6 | 2971 |
| 9 | 2956 |
| 26 | 2953 |
| 24 | 2940 |
| 28 | 2911 |
| 18 | 2904 |
| 33 | 2899 |
| 19 | 2893 |
| 21 | 2882 |
| 10 | 2870 |
| 1 | 2862 |
| 4 | 2861 |
| 30 | 2858 |
| 35 | 2855 |
| 32 | 2853 |
| 25 | 2853 |
| 14 | 2849 |
| 5 | 2847 |
| 16 | 2846 |
| 3 | 2846 |
| 7 | 2841 |
| 12 | 2832 |
| 23 | 2830 |
| 22 | 2826 |
| 8 | 2825 |
| 15 | 2823 |
| 17 | 2818 |
| 27 | 2815 |
| 2 | 2813 |
| 29 | 2794 |
| 20 | 2787 |
| 11 | 2767 |
| 34 | 2758 |
| 31 | 2736 |

Name: count, dtype: int64

Value counts for column 'Lead_Time':

Lead_Time

| | |
|----|------|
| 8 | 5428 |
| 1 | 5386 |
| 12 | 5368 |
| 5 | 5337 |
| 2 | 5318 |
| 17 | 5316 |

```
9      5263
14     5260
15     5258
6      5256
11     5255
19     5250
16     5242
13     5235
4      5217
10     5208
18     5184
3      5117
7      5102
Name: count, dtype: int64
```

```
Value counts for column 'Promotions_Applied':
Promotions_Applied
Yes      50261
No       49739
Name: count, dtype: int64
```

```
Value counts for column 'Return_Flag':
Return_Flag
Yes      50103
No       49897
Name: count, dtype: int64
```

```
Value counts for column 'User_Demographics':
User_Demographics
Individual      25117
Government     25067
Enterprise     24994
Small Business  24822
Name: count, dtype: int64
```

```
cat_cols = [col for col in df.columns if df[col].dtype == 'object' or
df[col].dtype.name == 'category']

for col in cat_cols:
    print(f"Value counts for column '{col}':")
```

```
print(df[col].value_counts())
print("\n" + "_"*40 + "\n")
```

Value counts for column 'Category':

```
Category
Circuit Breaker    20166
Switch              20024
Busway              20006
Power Meter         19945
Relay               19859
Name: count, dtype: int64
```

Value counts for column 'Technical_Specifications':

```
Technical_Specifications
{"voltage": "220V", "power": "50W"}    50074
{"voltage": "110V", "power": "30W"}    49926
Name: count, dtype: int64
```

Value counts for column 'Region':

```
Region
East      25177
South     25049
North     24887
West      24887
Name: count, dtype: int64
```

Value counts for column 'Fault_Flag':

```
Fault_Flag
Yes      50217
No       49783
Name: count, dtype: int64
```

Value counts for column 'Date':

```
Date
2020-01-01 00:00:00    1
2027-08-09 02:00:00    1
2027-08-10 00:00:00    1
2027-08-09 23:00:00    1
2027-08-09 22:00:00    1
..
2023-10-20 19:00:00    1
2023-10-20 18:00:00    1
```

```
2023-10-20 17:00:00    1
2023-10-20 16:00:00    1
2031-05-29 15:00:00    1
Name: count, Length: 100000, dtype: int64
```

```
Value counts for column 'Sensor_Readings':
Sensor_Readings
{"temperature": 40, "pressure": 1.5}    50084
{"temperature": 30, "pressure": 1.2}    49916
Name: count, dtype: int64
```

```
Value counts for column 'Promotions_Applied':
Promotions_Applied
Yes      50261
No       49739
Name: count, dtype: int64
```

```
Value counts for column 'Return_Flag':
Return_Flag
Yes      50103
No       49897
Name: count, dtype: int64
```

```
Value counts for column 'User_Demographics':
User_Demographics
Individual      25117
Government      25067
Enterprise      24994
Small Business  24822
Name: count, dtype: int64
```

EDA

```
num_features = df.select_dtypes(include = ['int64',
'float64']).dtypes.index

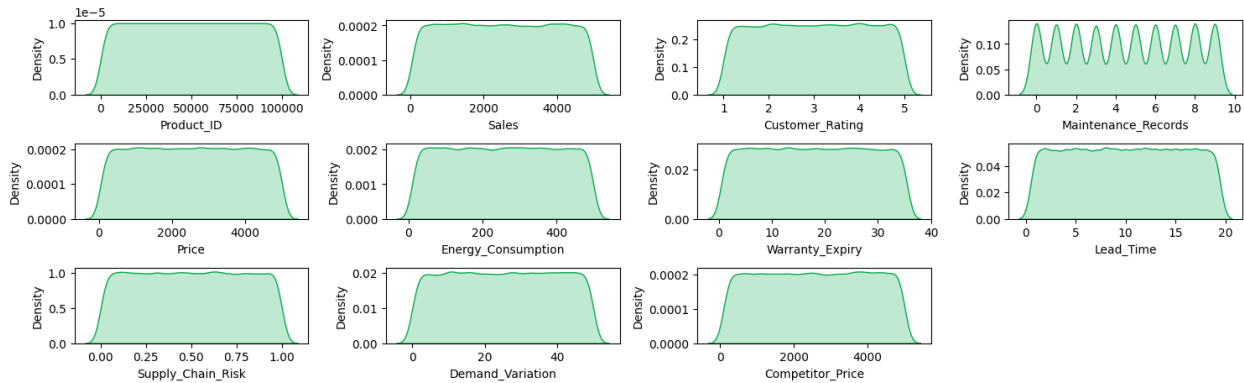
plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of
Features', fontweight='bold', fontsize=15, y=1)
```

```

for i in range(0,len(num_features)):
    plt.subplot(10,4,i+1)
    sns.kdeplot(x=df[num_features[i]],shade=True,color='#06a94d')
    plt.tight_layout()

```

Univariate Analysis of Features



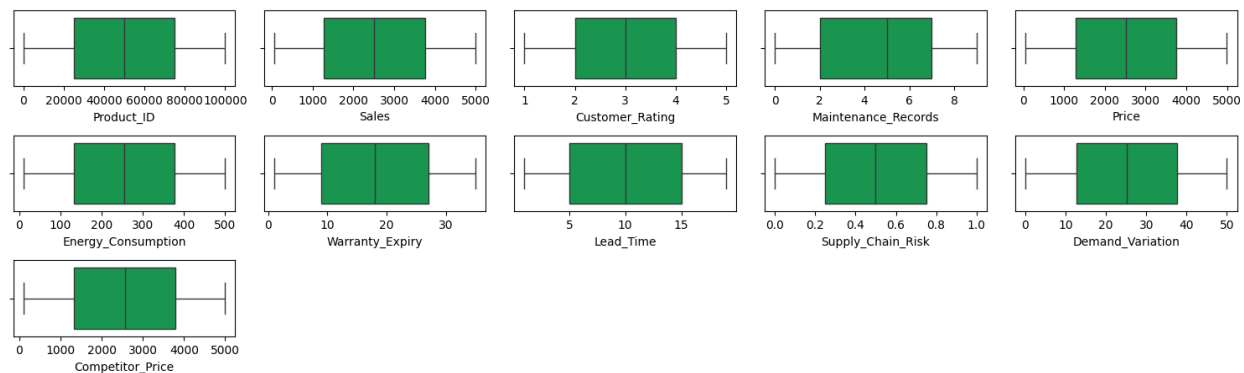
```

plt.figure(figsize = (15,15))
plt.suptitle('Univariate Analysis of
Features',fontweight='bold',fontsize=20,y=1)

for i in range(0,len(num_features)):
    plt.subplot(10,5,i+1)
    sns.boxplot(data=df,x=num_features[i],color='#06a94d')
    plt.xlabel(num_features[i])
    plt.tight_layout()

```

Univariate Analysis of Features



```

cat_features = df.select_dtypes(include='object').dtypes.index

plt.figure(figsize=(15,15))
plt.suptitle('Univariate Analysis of
Features',fontweight='bold',fontsize=15,y=1)

for i in range(0,len(cat_features)):

```

```
plt.subplot(10,4,i+1)
sns.countplot(x=df[cat_features[i]],color='#06a94d')
plt.tight_layout()
```

Feature Preprocessing

```
# Extract the values for Voltage and Power
df['Voltage'] = df['Technical_Specifications'].apply(lambda x: eval(x)
['voltage']).replace('V', '')
df['Power'] = df['Technical_Specifications'].apply(lambda x: eval(x)
['power']).replace('W', '')

# Extract temperature and pressure from the Sensor_Readings column
df['Temperature'] = df['Sensor_Readings'].apply(lambda x: eval(x)
['temperature'])
df['Pressure'] = df['Sensor_Readings'].apply(lambda x: eval(x)
['pressure'])

# Define the mapping function
def map_category(value):
    mapping = {
        "Circuit Breaker": "4",
        "Switch": "3",
        "Busway": "2",
        "Power Meter": "1",
        "Relay": "0"
    }
    return mapping.get(value, "Unknown") # Default to "Unknown" if
value not in mapping

# Apply the mapping function to the Category column
df['Category'] = df['Category'].apply(map_category)

# Define the mapping function
def map_fault_flag(value):
    mapping = {
        "Yes": 1,
        "No": 0
    }
    return mapping.get(value, -1) # Default to -1 if value not in
mapping

# Apply the mapping function to the Fault_Flag column
df['Fault_Flag'] = df['Fault_Flag'].apply(map_fault_flag)

# Define the mapping function
def map_return_flag(value):
    mapping = {
        "Yes": 1,
```

```

        "No": 0
    }
    return mapping.get(value, -1) # Default to -1 if value not in
mapping

# Apply the mapping function to the Return_Flag column
df['Return_Flag'] = df['Return_Flag'].apply(map_return_flag)

# Define the mapping function
def map_user_demographics(value):
    mapping = {
        "Individual": 0,
        "Government": 1,
        "Enterprise": 2,
        "Small Business": 3
    }
    return mapping.get(value, -1)

df['User_Demographics'] =
df['User_Demographics'].apply(map_user_demographics)

# Example Machine Learning Use Cases:

```

Product Fault Detection (Classification Problem)

- Objective: Predict whether a product will have a fault or not (Fault_Flag column).

```

df['Fault_Flag'].value_counts()

Fault_Flag
1    50217
0    49783
Name: count, dtype: int64

# Preprocessing
X = df[['Category', 'Maintenance_Records', 'Energy_Consumption',
'Warranty_Expiry', 'Voltage', 'Power']]
y = df['Fault_Flag'] # .apply(lambda x: 1 if x == 'Yes' else 0) #
Binary classification

# Imbalanced Data Handling
from imblearn.over_sampling import SMOTE

smote = SMOTE(random_state=42)
X_res,y_res = smote.fit_resample(X,y)

X_res,y_res = SMOTE().fit_resample(X,y)

from sklearn.model_selection import train_test_split

# Split into train and test sets

```

```

X_train, X_test, y_train, y_test = train_test_split(X_res, y_res,
test_size=0.20, random_state=42)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.fit_transform(X_test)

from sklearn.metrics import accuracy_score, classification_report

def get_best_model_report(models, X_train, y_train, X_test, y_test):
    """
    Train models, evaluate their accuracy, and display the
    classification report for the best model.

    Args:
        models (dict): A dictionary with model names as keys and model
        objects as values.
        X_train, y_train: Training features and labels.
        X_test, y_test: Testing features and labels.
    """
    best_model_name = None
    best_model = None
    best_accuracy = 0
    y_best_pred = None

    # Train, predict, and evaluate each model
    for model_name, model in models.items():
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        accuracy = accuracy_score(y_test, y_pred)
        print(f'Accuracy {model_name}: {accuracy:.4f}')

        if accuracy > best_accuracy:
            best_accuracy = accuracy
            best_model_name = model_name
            best_model = model
            y_best_pred = y_pred

    print(f"\nBest Model: {best_model_name} with Accuracy:
{best_accuracy:.4f}\n")
    print("Classification Report for Best Model:\n")
    print(classification_report(y_test, y_best_pred))

# Example usage
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier

```



```

from sklearn.ensemble import RandomForestClassifier,
GradientBoostingClassifier

# Define models
models = {
    "LogisticRegression": LogisticRegression(n),
    "KNeighborsClassifier": KNeighborsClassifier(n),
    "SVC": SVC(),
    "DecisionTreeClassifier": DecisionTreeClassifier(),
    "RandomForestClassifier": RandomForestClassifier(),
    "GradientBoostingClassifier": GradientBoostingClassifier()
}

# Call the function
get_best_model_report(models, X_train, y_train, X_test, y_test)

Accuracy LogisticRegression: 0.4949
Accuracy KNeighborsClassifier: 0.5012
Accuracy SVC: 0.4963
Accuracy DecisionTreeClassifier: 0.4965
Accuracy RandomForestClassifier: 0.4954
Accuracy GradientBoostingClassifier: 0.5003

Best Model: KNeighborsClassifier with Accuracy: 0.5012

Classification Report for Best Model:

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.50 | 0.50 | 0.50 | 10079 |
| 1 | 0.50 | 0.50 | 0.50 | 10008 |
| accuracy | | | 0.50 | 20087 |
| macro avg | 0.50 | 0.50 | 0.50 | 20087 |
| weighted avg | 0.50 | 0.50 | 0.50 | 20087 |

Sales Demand Forecasting (Regression Problem)

- Objective: Predict the sales of products based on historical data.

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

# Define the mapping function
def map_region(value):
    mapping = {
        "East": 3,
        "South": 2,
        "North": 1,
        "West": 0
    }

```

```

    }
    return mapping.get(value, -1)

df['Region'] = df['Region'].apply(map_region)

# Preprocessing
X = df[['Customer_Rating', 'Region', 'Energy_Consumption',
'Demand_Variation', 'Supply_Chain_Risk', 'Lead_Time', 'Price']]
y = df['Sales']

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Model training
regressor = RandomForestRegressor()
regressor.fit(X_train, y_train)

RandomForestRegressor()

# Predictions
y_pred = regressor.predict(X_test)

# Evaluation
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")

Mean Squared Error: 2108473.0428792466

import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score

# XGBoost model instance
xgb_model = xgb.XGBRegressor(objective='reg:squarederror',
random_state=42)

# Hyperparameter tuning with GridSearchCV
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.05, 0.1],
    'subsample': [0.8, 0.9, 1.0],
    'colsample_bytree': [0.8, 0.9, 1.0]
}

grid_search = GridSearchCV(estimator=xgb_model, param_grid=param_grid,
scoring='r2', cv=3, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

# Best parameters
best_params = grid_search.best_params_

```

```

print("Best parameters:", best_params)

# Train the model with the best parameters
best_xgb_model = xgb.XGBRegressor(**best_params,
objective='reg:squarederror', random_state=42)
best_xgb_model.fit(X_train, y_train)

# Evaluate on training and testing sets
y_train_pred = best_xgb_model.predict(X_train)
y_test_pred = best_xgb_model.predict(X_test)

# Calculate performance metrics
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

print(f"Training RMSE: {train_rmse}, R2: {train_r2}")
print(f"Testing RMSE: {test_rmse}, R2: {test_r2}")

# Check for underfitting and overfitting
if train_r2 > test_r2 + 0.1:
    print("The model might be overfitting.")
elif test_r2 > train_r2 + 0.1:
    print("The model might be underfitting.")
else:
    print("The model is performing well.")

Fitting 3 folds for each of 243 candidates, totalling 729 fits
Best parameters: {'colsample_bytree': 0.8, 'learning_rate': 0.01,
'max_depth': 3, 'n_estimators': 100, 'subsample': 1.0}
Training RMSE: 1428.672709539614, R2: 0.0010957717895507812
Testing RMSE: 1431.829846463929, R2: 3.427267074584961e-05
The model is performing well.

```

Predict Maintenance Needs (Classification Problem)

- Objective: Predict when maintenance will be required based on usage and sensor data.

```

from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

# Preprocessing
X = df[['Energy_Consumption', 'Supply_Chain_Risk', 'Warranty_Expiry',
'Temperature', 'Pressure']]
y = df['Maintenance_Records'].apply(lambda x: 1 if x > 0 else 0) #
Binary classification (maintenance or no maintenance)

```

```
# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

```
# Model training
gb_model = GradientBoostingClassifier()
gb_model.fit(X_train, y_train)
```

```
GradientBoostingClassifier()
```

```
# Predictions
y_pred = gb_model.predict(X_test)
```

```
# Evaluation
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
print(classification_report(y_test, y_pred))
```

```
Accuracy: 0.8993666666666666
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.00 | 0.00 | 0.00 | 3018 |
| 1 | 0.90 | 1.00 | 0.95 | 26982 |
| accuracy | | | 0.90 | 30000 |
| macro avg | 0.45 | 0.50 | 0.47 | 30000 |
| weighted avg | 0.81 | 0.90 | 0.85 | 30000 |

```
# Predict Maintenance Needs (Classification or Regression Problem)
# Objective: Predict when maintenance will be required based on usage
and sensor data.
```

```
# Model: Classification models for predicting maintenance or
regression models for predicting the time to next maintenance.
```

```
# Classification Problem
```

```
# The goal is to predict whether a product will require maintenance
(Maintenance_Need_Flag = 1) or not (Maintenance_Need_Flag = 0).
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
```

```
# Feature selection and preprocessing
```

```
X = df[['Energy_Consumption', 'Supply_Chain_Risk', 'Warranty_Expiry',
'Temperature', 'Pressure']]
```

```
y = df['Maintenance_Records'].apply(lambda x: 1 if x > 0 else 0) #
Binary classification (maintenance or no maintenance)
```

```
# Split the dataset
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

```

# Model training
clf_model = RandomForestClassifier(random_state=42)
clf_model.fit(X_train, y_train)

RandomForestClassifier(random_state=42)

# Predictions
y_pred = clf_model.predict(X_test)

# Evaluation
print(f"Classification Accuracy: {accuracy_score(y_test, y_pred)}")
print(classification_report(y_test, y_pred))

Classification Accuracy: 0.8873

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.10 | 0.01 | 0.03 | 3018 |
| 1 | 0.90 | 0.98 | 0.94 | 26982 |
| accuracy | | | 0.89 | 30000 |
| macro avg | 0.50 | 0.50 | 0.48 | 30000 |
| weighted avg | 0.82 | 0.89 | 0.85 | 30000 |

```

# Regression Problem
# The goal is to predict the time remaining before maintenance is
# required (in days or hours).

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error

# Normalize columns for calculation
df['Energy_Consumption_Norm'] = df['Energy_Consumption'] /
df['Energy_Consumption'].max()
df['Temperature_Norm'] = df['Temperature'] / df['Temperature'].max()
df['Pressure_Norm'] = df['Pressure'] / df['Pressure'].max()
# Define maintenance lead time (in days)
# The lower the combined normalized score, the shorter the maintenance
# lead time.
df['Maintenance_Lead_Time'] = 100 - (
    40 * df['Energy_Consumption_Norm'] +
    30 * df['Temperature_Norm'] +
    20 * df['Pressure_Norm']
)

# Add random noise to simulate variability
df['Maintenance_Lead_Time'] = df['Maintenance_Lead_Time'] +
np.random.normal(0, 5, len(df))

# Clip values to ensure no negative lead times
df['Maintenance_Lead_Time'] =

```

```

df['Maintenance_Lead_Time'].clip(lower=1)

# Drop temporary normalized columns
df.drop(columns=['Energy_Consumption_Norm', 'Temperature_Norm',
                 'Pressure_Norm'], inplace=True)

# Feature selection and preprocessing
X = df[['Maintenance_Records', 'Energy_Consumption',
        'Supply_Chain_Risk', 'Warranty_Expiry', 'Temperature', 'Pressure']]
y = df['Maintenance_Lead_Time'] # Numeric column indicating time to
maintenance

# Split the dataset
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

# Model training
reg_model = RandomForestRegressor(random_state=42)
reg_model.fit(X_train, y_train)

RandomForestRegressor(random_state=42)

# Predictions
y_pred = reg_model.predict(X_test)

# Evaluation
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")

Mean Squared Error: 27.291767909916082
Mean Absolute Error: 4.170469035139971

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score

# Random Forest model instance
rf_model = RandomForestRegressor(random_state=42)

# Hyperparameter tuning with GridSearchCV
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}

grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid,

```

```

scoring='r2', cv=3, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)

# Best parameters
best_params = grid_search.best_params_
print("Best parameters:", best_params)

# Train the model with the best parameters
best_rf_model = RandomForestRegressor(**best_params, random_state=42)
best_rf_model.fit(X_train, y_train)

# Evaluate on training and testing sets
y_train_pred = best_rf_model.predict(X_train)
y_test_pred = best_rf_model.predict(X_test)

# Calculate performance metrics
train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

train_r2 = r2_score(y_train, y_train_pred)
test_r2 = r2_score(y_test, y_test_pred)

print(f"Training RMSE: {train_rmse:.4f}, R2: {train_r2:.4f}")
print(f"Testing RMSE: {test_rmse:.4f}, R2: {test_r2:.4f}")

# Check for underfitting and overfitting
if train_r2 > test_r2 + 0.1:
    print("The model might be overfitting.")
elif test_r2 > train_r2 + 0.1:
    print("The model might be underfitting.")
else:
    print("The model is performing well.")

Fitting 3 folds for each of 324 candidates, totalling 972 fits
Best parameters: {'max_depth': 20, 'max_features': 'sqrt',
'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 150}
Training RMSE: 4.0193, R2: 0.9135
Testing RMSE: 5.0704, R2: 0.8613
The model is performing well.

```

Product Return Prediction (Classification Problem)

- Objective: Predict whether a product will be returned (Return_Flag).

```

df['Return_Flag'].value_counts()

Return_Flag
1    50103
0    49897
Name: count, dtype: int64

```

```

# Preprocessing
X = df[['Customer_Rating', 'Price', 'Sales', 'Warranty_Expiry',
'Category', 'User_Demographics']]
y = df['Return_Flag'] # .apply(lambda x: 1 if x == 'Yes' else 0)

# Imbalanced Data Handling
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split

smote = SMOTE(random_state=42)
X_res,y_res = smote.fit_resample(X,y)

X_res,y_res = SMOTE().fit_resample(X,y)

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res,
test_size=0.20, random_state=42)

# Model training
model = RandomForestClassifier()
model.fit(X_train, y_train)

RandomForestClassifier()

# Predictions
y_pred = model.predict(X_test)

# Evaluation
print(f"Accuracy: {accuracy_score(y_test, y_pred)}")
print(classification_report(y_test, y_pred))

Accuracy: 0.5045903602434887

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.50 | 0.53 | 0.52 | 9966 |
| 1 | 0.51 | 0.48 | 0.49 | 10076 |
| accuracy | | | 0.50 | 20042 |
| macro avg | 0.50 | 0.50 | 0.50 | 20042 |
| weighted avg | 0.50 | 0.50 | 0.50 | 20042 |

```

import pickle

pickle.dump(model,open('product_return_predictor.pkl','wb'))

# model = pickle.load(open('product_return_predictor.pkl','rb'))

# Assuming `model` is already trained and loaded

print("Product Return Prediction")

```



```

# Input data
customer_rating = float(input("Enter the Customer Rating: "))
price = float(input("Enter the Price: "))
sales = int(input("Enter the Sales: "))
warranty_expiry = int(input("Enter the Warranty Expiry: "))
category = int(input("Enter the Category: "))
user_demographics = int(input("Enter the User Demographics: "))

# Create a feature array
input_point = np.array([[customer_rating, price, sales,
warranty_expiry, category, user_demographics]])

# Make a prediction
prediction = model.predict(input_point)

# Check the prediction and print the result
if prediction[0] == 1:
    print("Product Return")
else: # No need to check for `prediction[0] == 0` since it's implied
    print("Product Not Returned")

```

```

Product Return Prediction
Enter the Customer Rating: 1.80
Enter the Price: 2271.00
Enter the Sales: 4713
Enter the Warranty Expiry: 7
Enter the Category: 3
Enter the User Demographics: 3
Product Not Returned

```

```

# X['Category'] = X['Category'].astype('int')

```

```

# Model training
x_model = XGBClassifier()
x_model.fit(X_train, y_train)

```

```

# Predictions
x_pred = model.predict(X_test)

```

```

# Evaluation
print(f"Accuracy: {accuracy_score(y_test, x_pred)}")
print(classification_report(y_test, x_pred))

```

```

Accuracy: 0.5040914080431095

```

| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 0 | 0.50 | 0.53 | 0.52 | 9966 |
| 1 | 0.51 | 0.48 | 0.49 | 10076 |
| accuracy | | | 0.50 | 20042 |
| macro avg | 0.50 | 0.50 | 0.50 | 20042 |

| | | | | |
|--------------|------|------|------|-------|
| weighted avg | 0.50 | 0.50 | 0.50 | 20042 |
|--------------|------|------|------|-------|

Code for XGBoost Classification with Hyperparameter Tuning

```
import xgboost as xgb
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
```

XGBoost classifier instance

```
xgb_clf = xgb.XGBClassifier(objective='binary:logistic',
use_label_encoder=False, eval_metric='logloss', random_state=42)
```

Hyperparameter tuning with GridSearchCV

```
param_grid = {
    'n_estimators': [50, 100, 150],
    'max_depth': [3, 5, 7],
    'learning_rate': [0.01, 0.05, 0.1],
    'subsample': [0.8, 0.9, 1.0],
    'colsample_bytree': [0.8, 0.9, 1.0],
    'gamma': [0, 0.1, 0.2],
    'reg_alpha': [0, 0.01, 0.1],
    'reg_lambda': [1, 1.5, 2]
}
```

```
grid_search = GridSearchCV(estimator=xgb_clf, param_grid=param_grid,
scoring='accuracy', error_score='raise', cv=3, verbose=1, n_jobs=-1)
grid_search.fit(X_train, y_train)
```

Best parameters

```
best_params = grid_search.best_params_
print("Best parameters:", best_params)
```

Train the model with the best parameters

```
best_xgb_model = xgb.XGBClassifier(**best_params,
objective='binary:logistic', use_label_encoder=False,
eval_metric='logloss', random_state=42, n_jobs=1)
best_xgb_model.fit(X_train, y_train)
```

Evaluate on training and testing sets

```
y_train_pred = best_xgb_model.predict(X_train)
y_test_pred = best_xgb_model.predict(X_test)
```

Calculate performance metrics

```
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)
```

```
print(f"Training Accuracy: {train_accuracy:.4f}")
```

```
print(f"Testing Accuracy: {test_accuracy:.4f}")
```

```

# Display classification report and confusion matrix
print("\nClassification Report:\n", classification_report(y_test,
y_test_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_test_pred))

# Check for overfitting or underfitting
if train_accuracy > test_accuracy + 0.1:
    print("The model might be overfitting.")
elif test_accuracy > train_accuracy + 0.1:
    print("The model might be underfitting.")
else:
    print("The model is performing well.")

```

Predicting Warranty Expiry Risk (Regression)

- Predict the time remaining on the warranty (Warranty_Expiry).

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Features and target
X = df[['Maintenance_Records', 'Voltage', 'Power', 'Temperature',
'Pressure']] # Maintenance_Lead_Time
y = df['Warranty_Expiry'] # .apply(lambda x: 1 if x == 'Yes' else 0)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Model training
regressor = RandomForestRegressor(random_state=42)
regressor.fit(X_train, y_train)

# Predictions and evaluation
y_pred = regressor.predict(X_test)
print("Mean Absolute Error:", mean_absolute_error(y_test, y_pred))
print("Mean Squared Error:", mean_squared_error(y_test, y_pred))
print("R2 Score:", r2_score(y_test, y_pred))

```

```

Mean Absolute Error: 8.731453395728098
Mean Squared Error: 101.68509626570433
R2 Score: -0.0004927865631072503

```

Collaborative Filtering (User-Product Recommendation)

- Recommend products to users based on historical sales and customer ratings.

```

from sklearn.metrics.pairwise import cosine_similarity
from sklearn.preprocessing import MinMaxScaler

```

```

# Load dataset
df_2 = pd.read_csv("schneider_electric_dataset.csv")

# Create a user-product matrix
user_product_matrix = df_2.pivot_table(
    index='Region', columns='Product_ID', values='Customer_Rating',
    aggfunc='mean'
).fillna(0)

# Compute similarity
user_similarity = cosine_similarity(user_product_matrix)
user_similarity_df = pd.DataFrame(user_similarity,
    index=user_product_matrix.index, columns=user_product_matrix.index)

# Recommend products for a specific user
def recommend_products(user, top_n=5):
    similar_users =
    user_similarity_df[user].sort_values(ascending=False)[1:top_n + 1]
    recommended_products =
    df_2[df_2['Region'].isin(similar_users.index)]['Product_ID'].unique()
    return recommended_products

# Example recommendation for a user
print(recommend_products('North', top_n=5))

[    2     3     4 ... 99997 99998 99999]

```

Content-Based Filtering (Product Similarity Recommendation)

- Recommend similar products based on Technical_Specifications, Category, and Sensor_Readings.

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.neighbors import NearestNeighbors

# Combine features for product similarity
df_2['Combined_Features'] = df_2['Category'] + " " +
df_2['Technical_Specifications'] + " " + df_2['Sensor_Readings']

# TF-IDF Vectorizer
tfidf = TfidfVectorizer(stop_words='english')
tfidf_matrix = tfidf.fit_transform(df_2['Combined_Features'])

# Nearest Neighbors model
nn = NearestNeighbors(metric='cosine', algorithm='brute',
    n_neighbors=6)
nn.fit(tfidf_matrix)

NearestNeighbors(algorithm='brute', metric='cosine', n_neighbors=6)

```

```

# Product recommendation function
def get_recommendations(product_id, top_n=5):
    product_idx = df_2[df_2['Product_ID'] == product_id].index[0]
    distances, indices = nn.kneighbors(tfidf_matrix[product_idx],
n_neighbors=top_n + 1)
    product_indices = indices.flatten()[1:] # Exclude the first one
(itself)
    return df_2['Product_ID'].iloc[product_indices].tolist()

# Example recommendation
print(get_recommendations(227, top_n=5))

[79689, 79719, 46712, 46716, 79738]

```

Region-Based Recommendation

- Recommend top-selling products in a specific Region.

```

# Get top products by region
def recommend_top_products(region, top_n=5):
    regional_sales = df_2[df_2['Region'] ==
region].groupby('Product_ID')['Sales'].sum()
    top_products =
regional_sales.sort_values(ascending=False).head(top_n).index.tolist()
    return top_products

# Example for North Region
print(recommend_top_products('East', top_n=5))

[31433, 97295, 95349, 88942, 12037]

```

Price-Sensitive Recommendation

- Recommend products within a similar price range.

```

# Recommend products in a price range
def recommend_by_price(product_id, price_range=100):
    product_price = df_2[df_2['Product_ID'] == product_id]
['Price'].values[0]
    recommended_products = df[
        (df_2['Price'] >= product_price - price_range) &
(df_2['Price'] <= product_price + price_range)
    ]['Product_ID'].tolist()
    return recommended_products

# Example recommendation
print(recommend_by_price(727, price_range=5))

[62, 607, 648, 727, 2212, 3080, 4357, 4518, 5243, 5589, 5690, 5863,
6288, 6447, 6736, 6868, 7404, 7594, 8127, 8584, 9223, 9312, 10096,
10606, 10976, 11092, 11279, 11387, 11588, 12540, 13201, 13674, 13787,

```

```

14022, 14077, 14516, 14573, 15363, 15562, 15580, 16061, 16132, 17202,
17219, 17875, 18260, 18717, 18904, 19535, 19739, 19771, 19867, 19983,
20464, 20572, 22235, 22244, 22309, 23242, 23746, 24056, 26336, 26499,
26777, 28004, 28602, 28806, 29270, 29349, 29904, 30033, 30922, 31044,
32630, 33722, 33983, 34236, 34432, 35398, 35623, 36697, 36993, 37289,
38195, 38360, 38417, 38544, 38554, 38628, 38949, 39222, 39288, 39383,
39549, 40104, 40224, 40722, 41476, 41892, 43090, 43333, 44375, 44812,
45587, 48165, 51023, 51205, 51290, 51423, 51658, 51868, 52199, 54147,
54870, 55439, 55675, 55751, 55977, 56134, 56154, 56779, 56986, 57636,
58520, 58746, 58748, 58971, 59439, 59771, 60495, 60751, 60819, 61597,
62397, 62426, 63429, 63795, 63937, 64227, 65080, 65621, 66605, 67187,
68621, 69167, 69593, 69673, 69895, 70788, 70863, 72609, 73307, 74435,
74663, 75782, 76808, 78202, 79026, 79914, 80025, 80193, 80620, 80662,
81069, 81430, 81914, 82166, 82324, 83826, 83871, 84049, 84483, 84491,
84849, 84966, 85158, 85212, 85681, 87251, 87272, 87677, 89491, 90478,
90678, 91693, 92265, 92415, 93204, 93823, 94669, 94997, 97111, 97536,
97772, 98454, 98559, 98663, 98730, 99788]

```

Customer Behavior Recommendation

- Recommend products based on customer demographics (User_Demographics).

```

# Recommend products for a specific user demographic
def recommend_by_demographics(demographic, top_n=5):
    demo_products = df_2[df_2['User_Demographics'] ==
demographic].groupby('Product_ID')['Sales'].sum()
    top_products =
demo_products.sort_values(ascending=False).head(top_n).index.tolist()
    return top_products

# Example for Small Business demographic
print(recommend_by_demographics('Small Business', top_n=5))

[95349, 94738, 37180, 77266, 97295]

```

Competitor Price-Based Recommendation

- Recommend products that are competitively priced.

```

# Recommend products with competitive pricing
def recommend_competitive_products(product_id, price_diff=50):
    product_price = df_2[df_2['Product_ID'] == product_id]
['Price'].values[0]
    recommended_products = df[
        abs(df_2['Price'] - df_2['Competitor_Price']) <= price_diff
    ]['Product_ID'].tolist()
    return recommended_products

# Example recommendation
print(recommend_competitive_products(227, price_diff=5))

```

```
[234, 932, 1207, 1350, 2966, 3302, 3750, 4259, 4804, 5237, 5329, 5784,
5858, 6473, 6716, 6939, 7124, 8708, 9528, 9793, 10058, 10606, 10898,
12142, 12298, 12643, 14023, 14530, 14713, 15945, 16253, 16360, 16491,
16972, 17402, 18726, 19179, 20267, 20353, 21165, 21334, 22279, 23514,
23541, 23841, 24511, 24645, 24920, 25159, 26317, 26323, 26688, 27304,
28634, 29036, 29521, 30407, 32719, 33141, 33556, 33925, 34587, 35366,
35389, 35441, 36426, 37039, 37366, 37448, 37723, 38161, 38261, 38325,
38513, 39562, 39880, 40327, 40601, 40650, 40943, 42055, 42458, 42615,
43271, 43332, 43586, 43662, 44913, 45106, 46848, 47350, 47513, 48043,
48327, 48374, 48687, 49203, 49284, 49519, 50000, 50082, 50105, 50175,
50235, 51021, 51091, 51627, 51953, 52107, 52755, 52838, 53037, 53115,
53269, 53544, 53948, 54134, 54539, 54571, 54913, 55484, 56014, 58071,
58108, 58153, 60111, 60244, 60394, 60875, 60919, 61214, 62059, 62282,
63112, 63795, 64259, 65501, 67218, 67281, 68200, 69109, 69563, 70219,
71636, 71680, 71842, 73305, 73328, 73332, 74249, 75527, 75825, 76198,
76866, 77416, 79182, 79775, 79890, 79921, 80509, 81360, 81589, 82168,
83166, 83258, 83710, 83768, 84844, 84907, 85053, 85418, 85615, 85768,
86047, 86430, 86703, 86982, 87142, 87320, 87439, 87506, 88364, 88852,
88864, 89000, 89323, 89791, 90340, 90962, 91514, 91576, 92093, 92313,
93217, 93781, 94039, 94668, 95304, 96653, 96690, 97545, 97731, 98142,
99121, 99229, 99458]
```

Energy Efficiency Recommendation

- Recommend products with lower energy consumption.

```
# Recommend energy-efficient products
def recommend_energy_efficient(top_n=5):
    efficient_products =
df_2.sort_values(by='Energy_Consumption').head(top_n)
['Product_ID'].tolist()
    return efficient_products

# Example recommendation
print(recommend_energy_efficient(top_n=5))

[7607, 68749, 24660, 61109, 82260]

# More Advance Working Sonn..!
# Notebook Project By : PRASAD JADHAV (ML-ENG)
# LinkedIn: linkedin.com/in/prasadmjadhav2 | Github:
github.com/prasadmjadhav2 | Mail: prasadmjadhav6161@gmail.com
```

Thank You!